



Spark: Resilient Distributed Datasets as Workflow System

빅데이터분석
천세진

Where is MapReduce Inefficient?

DFS → Map → LocalFS → Network → Reduce → DFS → Map → ...



Where is MapReduce Inefficient?

- Long pipelines sharing data
- Interactive applications
- Streaming applications
- Iterative algorithms (optimization problems)

DFS → Map → LocalFS → Network → Reduce → DFS → Map → ...




Spark의 Big Idea

- Resilient Distributed Datasets(RDDs)
Read-only partitioned collection of records (like a DFS)
- But, 어떻게 데이터가 생성되었는지에 대한 레코드를 가짐
 - Combination of *transformations* from other dataset(s).



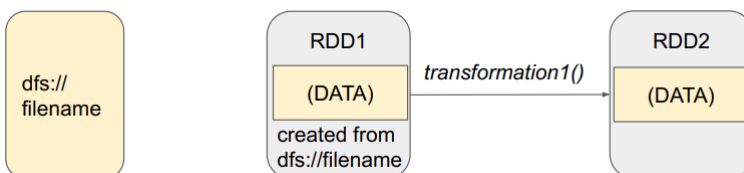
Spark의 Big Idea

- Resilient Distributed Datasets(RDDs)
Read-only partitioned collection of records (like a DFS)
- But, 어떻게 데이터가 생성되었는지에 대한 레코드를 가짐
 - Combination of *transformations* from other dataset(s).



Spark의 Big Idea

- Resilient Distributed Datasets(RDDs)
Read-only partitioned collection of records (like a DFS)
- But, 어떻게 데이터가 생성되었는지에 대한 레코드를 가짐
 - Combination of *transformations* from other dataset(s).



Spark의 Big Idea

- Resilient Distributed Datasets(RDDs)
Read-only partitioned collection of records (like a DFS)
- But, 어떻게 데이터가 생성되었는지에 대한 레코드를 가짐
 - Combination of *transformations* from other dataset(s).




Spark의 Big Idea

- Resilient Distributed Datasets(RDDs)
Read-only partitioned collection of records (like a DFS)
- But, 어떻게 데이터가 생성되었는지에 대한 레코드를 가짐
 - Combination of *transformations* from other dataset(s).
- Faster communication and I/O
 - On-the-fly 형태로 데이터셋을 rebuilding이 가능함
 - disk에 중간 결과(Intermediate datasets)가 저장되지 않음
 - Only in-memory




Spark의 Big Idea

- Resilient Distributed Datasets(RDDs)
Read-only partitioned collection of records (like a DFS)
- But, 어떻게 데이터가 생성되었는지에 대한 레코드를 가짐
 - Combination of *transformations* from **other dataset(s)**.



Spark의 Big Idea

- Resilient Distributed Datasets(RDDs)
Read-only partitioned collection of records (like a DFS)
- But, 어떻게 데이터가 생성되었는지에 대한 레코드를 가짐
 - Combination of *transformations* from other dataset(s).



Spark의 Big Idea

■ Resilient Distributed Datasets(RDDs)

Read-only partitioned collection of records (like a DFS)

■ But, 어떻게 데이터가 생성되었는지에 대한 레코드를 가짐

- Combination of *transformations* from other dataset(s).



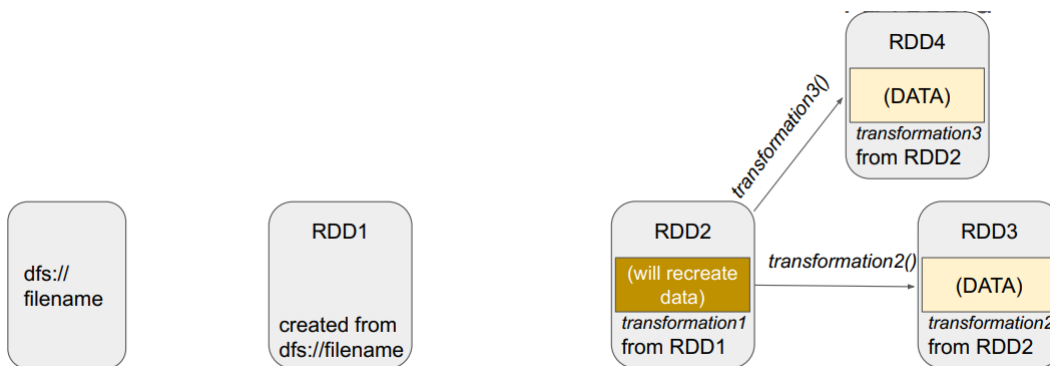
Spark의 Big Idea

■ Resilient Distributed Datasets(RDDs)

Read-only partitioned collection of records (like a DFS)

■ But, 어떻게 데이터가 생성되었는지에 대한 레코드를 가짐

- Combination of *transformations* from other dataset(s).



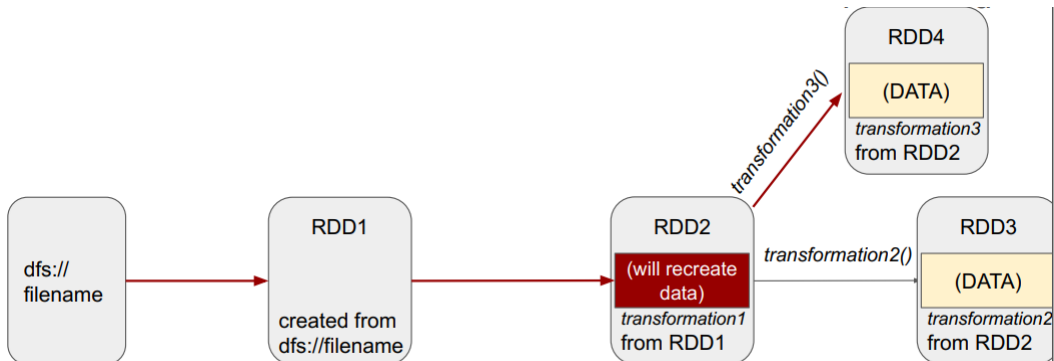
Spark의 Big Idea

■ Resilient Distributed Datasets(RDDs)

Read-only partitioned collection of records (like a DFS)

■ But, 어떻게 데이터가 생성되었는지에 대한 레코드를 가짐

- Combination of *transformations* from other dataset(s).



Transformations: RDD to RDD

| | | |
|-----------------|---|---|
| Transformations | $map(f : T \Rightarrow U)$ | : $RDD[T] \Rightarrow RDD[U]$ |
| | $filter(f : T \Rightarrow Bool)$ | : $RDD[T] \Rightarrow RDD[T]$ |
| | $flatMap(f : T \Rightarrow Seq[U])$ | : $RDD[T] \Rightarrow RDD[U]$ |
| | $sample(fraction : Float)$ | : $RDD[T] \Rightarrow RDD[T]$ (Deterministic sampling) |
| | $groupByKey()$ | : $RDD[(K, V)] \Rightarrow RDD[(K, Seq[V])]$ |
| | $reduceByKey(f : (V, V) \Rightarrow V)$ | : $RDD[(K, V)] \Rightarrow RDD[(K, V)]$ |
| | $union()$ | : $(RDD[T], RDD[T]) \Rightarrow RDD[T]$ |
| | $join()$ | : $(RDD[(K, V)], RDD[(K, W)]) \Rightarrow RDD[(K, (V, W))]$ |
| | $cogroup()$ | : $(RDD[(K, V)], RDD[(K, W)]) \Rightarrow RDD[(K, (Seq[V], Seq[W]))]$ |
| | $crossProduct()$ | : $(RDD[T], RDD[U]) \Rightarrow RDD[(T, U)]$ |
| | $mapValues(f : V \Rightarrow W)$ | : $RDD[(K, V)] \Rightarrow RDD[(K, W)]$ (Preserves partitioning) |
| | $sort(c : Comparator[K])$ | : $RDD[(K, V)] \Rightarrow RDD[(K, V)]$ |
| | $partitionBy(p : Partitioner[K])$ | : $RDD[(K, V)] \Rightarrow RDD[(K, V)]$ |

Table 2: Transformations and actions available on RDDs in Spark. Seq[T] denotes a sequence of elements of type T.

Transformations: RDD to Value Object, or Storage

eager 와 lazy 실행의 설명

Actions

$count() : RDD[T] \Rightarrow Long$
 $collect() : RDD[T] \Rightarrow Seq[T]$
 $reduce(f : (T, T) \Rightarrow T) : RDD[T] \Rightarrow T$
 $lookup(k : K) : RDD[(K, V)] \Rightarrow Seq[V]$ (On hash/range partitioned RDDs)
 $save(path : String) : \text{Outputs RDD to a storage system, e.g., HDFS}$



Current Transformation and Actions

- filter, map, flatMap, reduceByKey, groupByKey


- collect, count, take



Example

Count errors in a log file:

TYPE MESSAGE TIME



Matei Zaharia, Mosharaf Chowdhury, Tathagata Das, Ankur Dave, Justin Ma, Murphy McCauley, Michael J. Franklin, Scott Shenker, Ion Stoica. "Resilient Distributed Datasets: A Fault-Tolerant Abstraction for In-Memory Cluster Computing.". *NSDI 2012*. April 2012.



Example


Count errors in a log file:

TYPE MESSAGE TIME

Pseudocode:

```

lines = sc.textFile("dfs:...")
errors =
    lines.filter(_.startswith("ERROR"))
errors.count
  
```



Matei Zaharia, Mosharaf Chowdhury, Tathagata Das, Ankur Dave, Justin Ma, Murphy McCauley, Michael J. Franklin, Scott Shenker, Ion Stoica. "Resilient Distributed Datasets: A Fault-Tolerant Abstraction for In-Memory Cluster Computing.". *NSDI 2012*. April 2012.




Example

Collect times of hdfs-related errors

TYPE MESSAGE TIME

Pseudocode:

```
lines = sc.textFile("dfs:...")
errors =
    lines.filter(_.startswith("ERROR"))
errors.persist
errors.count
...
```



Matei Zaharia, Mosharaf Chowdhury, Tathagata Das, Ankur Dave, Justin Ma, Murphy McCauley, Michael J. Franklin, Scott Shenker, Ion Stoica. "Resilient Distributed Datasets: A Fault-Tolerant Abstraction for In-Memory Cluster Computing.", NSDI 2012, April 2012.



Example

Collect times of hdfs-related errors

TYPE MESSAGE TIME

Pseudocode:

```
lines = sc.textFile("dfs:...")
errors =
    lines.filter(_.startswith("ERROR"))
errors.persist
errors.count
...
```

Persistence

Can specify that an RDD "persists" in memory so other queries can use it.

Can specify a priority for persistence; lower priority => moves to disk, if needed, earlier

Matei Zaharia, Mosharaf Chowdhury, Tathagata Das, Ankur Dave, Justin Ma, Murphy McCauley, Michael J. Franklin, Scott Shenker, Ion Stoica. "Resilient Distributed Datasets: A Fault-Tolerant Abstraction for In-Memory Cluster Computing.", NSDI 2012, April 2012.



Example

Collect times of hdfs-related errors

TYPE MESSAGE TIME

Pseudocode:

```
lines = sc.textFile("dfs:...")
errors =
    lines.filter(_.startswith("ERROR"))
errors.persist
errors.count
...
```

Persistence

Can specify that an RDD "persists" in memory so other queries can use it.

Can specify a priority for persistence; lower priority => moves to disk, if needed, earlier

[parameters for persist](#)

Matei Zaharia, Mosharaf Chowdhury, Tathagata Das, Ankur Dave, Justin Ma, Murphy McCauley, Michael J. Franklin, Scott Shenker, Ion Stoica. "Resilient Distributed Datasets: A Fault-Tolerant Abstraction for In-Memory Cluster Computing." *NSDI 2012*, April 2012.




Example

Collect times of hdfs-related errors

TYPE MESSAGE TIME

Pseudocode:

```
lines = sc.textFile("dfs:...")
errors =
    lines.filter(_.startswith("ERROR"))
errors.persist
errors.count
errors.filter(_.contains("HDFS"))
...
```



Matei Zaharia, Mosharaf Chowdhury, Tathagata Das, Ankur Dave, Justin Ma, Murphy McCauley, Michael J. Franklin, Scott Shenker, Ion Stoica. "Resilient Distributed Datasets: A Fault-Tolerant Abstraction for In-Memory Cluster Computing." *NSDI 2012*, April 2012.




Example

Collect times of hdfs-related errors

TYPE MESSAGE TIME

Pseudocode:

```
lines = sc.textFile("dfs:...")
errors =
  lines.filter(_.startswith("ERROR"))
errors.persist
errors.count
errors.filter(_.contains("HDFS"))
  .map(_split('\t')(3))
  .collect()
```



Matei Zaharia, Mosharaf Chowdhury, Tathagata Das, Ankur Dave, Justin Ma, Murphy McCauley, Michael J. Franklin, Scott Shenker, Ion Stoica. "Resilient Distributed Datasets: A Fault-Tolerant Abstraction for In-Memory Cluster Computing." NSDI 2012, April 2012.



Example

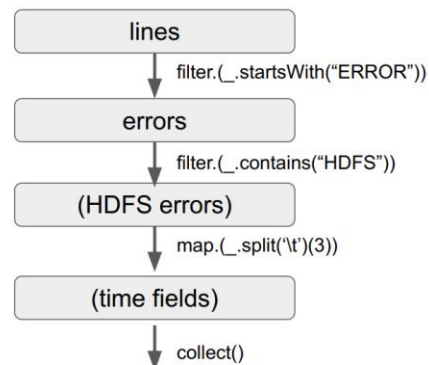
Collect times of hdfs-related errors

TYPE MESSAGE TIME

Pseudocode:

```
lines = sc.textFile("dfs:...")
errors =
  lines.filter(_.startswith("ERROR"))
errors.persist
errors.count
errors.filter(_.contains("HDFS"))
  .map(_split('\t')(3))
  .collect()
```

Functional Programming



Matei Zaharia, Mosharaf Chowdhury, Tathagata Das, Ankur Dave, Justin Ma, Murphy McCauley, Michael J. Franklin, Scott Shenker, Ion Stoica. "Resilient Distributed Datasets: A Fault-Tolerant Abstraction for In-Memory Cluster Computing." NSDI 2012, April 2012.



Example

Collect times of hdfs-related errors

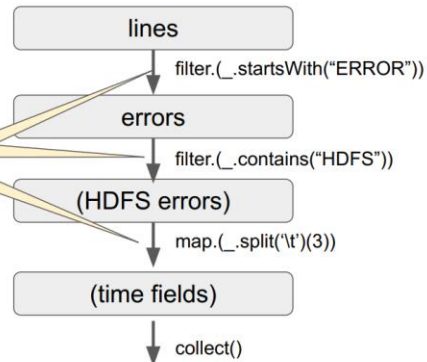
TYPE MESSAGE TIME

Pseudocode:

```
lines = sc.textFile("dfs:...")
errors =
  lines.filter(_.startswith("ERROR"))
errors.persist
errors.count
errors.filter(_.contains("HDFS"))
  .map(_split('\t')(3))
  .collect()
```

Functional Programming

"lineage"



• Matei Zaharia, Mosharaf Chowdhury, Tathagata Das, Ankur Dave, Justin Ma, Murphy McCauley, Michael J. Franklin, Scott Shenker, Ion Stoica: "Resilient Distributed Datasets: A Fault-Tolerant Abstraction for In-Memory Cluster Computing.", *NSDI 2012*, April 2012.

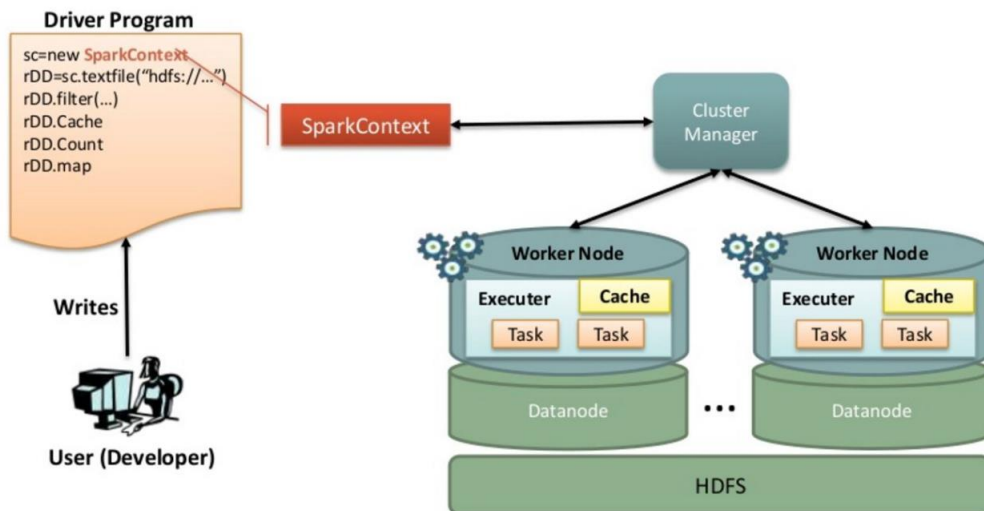


Workflow System의 장점

- More efficient failure recovery
- More efficient grouping of tasks and scheduling
- Integration of programming language features:
 - Loops (not a "cyclic" workflow system)
 - Function libraries



The Spark Programming Model



Gupta, Manish. Lightening Fast Big Data Analytics using Apache Spark. *UniCom 2014*.

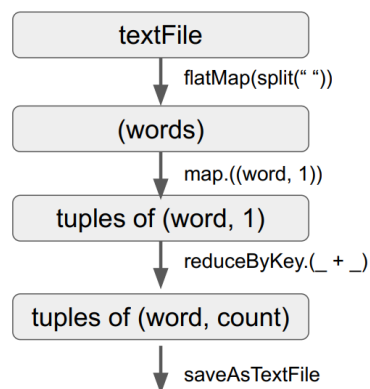


The Spark Programming Model

Word Count

Scala:

```
val textFile =
  sc.textFile("hdfs://...")
val counts = textFile
  .flatMap(line => line.split(" "))
  .map(word => (word, 1))
  .reduceByKey(_ + _)
counts.saveAsTextFile("hdfs://...")
```



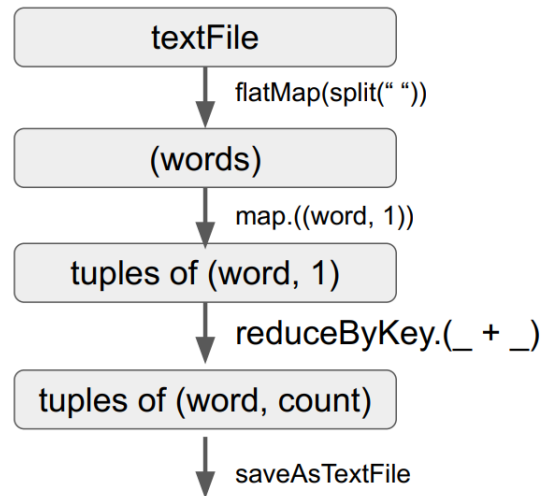
Apache Spark Examples
<http://spark.apache.org/examples.html>



Word Count

Python:

```
textFile = sc.textFile("hdfs://...")
counts = textFile
    .flatMap(lambda line: line.split(" "))
    .map(lambda word: (word, 1))
    .reduceByKey(lambda a, b: a + b)
counts.saveAsTextFile("hdfs://...")
```



Apache Spark Examples
<http://spark.apache.org/examples.html>

Lazy Evaluation

Spark waits to load data and execute transformations until necessary -- *lazy*
 Spark tries to complete actions as immediately as possible -- *eager*

Why?

- Only executes what is necessary to achieve action.
- Can optimize the complete chain of operations to reduce communication



Lazy Evaluation

Spark waits to *load data* and *execute transformations* until necessary -- **lazy**

Spark tries to complete actions as quickly as possible -- **eager**

Why?

- Only executes what is necessary to achieve action.
- Can optimize the complete *chain of operations* to reduce communication

e.g.

rdd(T) -> rdd(T)

```
rdd.map(lambda r: r[1]*r[3]).take(5) #only executes map for five records
rdd.filter(lambda r: "ERROR" in r[0]).map(lambda r: r[1]*r[3])
#only passes through the data once
```



Broadcast Variables

Read-only objects can be shared across all nodes.

Broadcast variable is a wrapper: access object with `.value`

Python:

```
filterWords = ['one', 'two', 'three', 'four', ...]
fwBC = sc.broadcast(set(filterWords))
```



Broadcast Variables

Read-only objects can be shared across all nodes.

Broadcast variable is a wrapper: access object with `.value`

```
Python:
filterWords = ['one', 'two', 'three', 'four', ...]
fwBC = sc.broadcast(set(filterWords))

textFile = sc.textFile("hdfs:...")
counts = textFile
    .map(lambda line: line.split(" "))
    .filter(lambda words: len(set(words)) and word in fwBC.value) > 0)
    .flatMap(lambda word: (word, 1))
    .reduceByKey(lambda a, b: a + b)
counts.saveAsTextFile("hdfs:...")
```



Accumulators

Write-only objects that keep a running aggregation

Default Accumulator assumes sum function

```
initialValue = 0
sumAcc = sc.accumulator(initialValue)
rdd.foreach(lambda i: sumAcc.add(i))
print(sumAcc.value)
```



Accumulators

Write-only objects that keep a running aggregation

Default Accumulator assumes sum function

Custom Accumulator: Inherit (AccumulatorParam) as class and override methods

```

initialValue = 0
sumAcc = sc.accumulator(initialValue)
rdd.foreach(lambda i: sumAcc.add(i))
print(minAcc.value)

class MinAccum(AccumulatorParam):
    def zero(self, zeroValue = np.inf):#overwrite this
        return zeroValue
    def addInPlace(self, v1, v2):#overwrite this
        return min(v1, v2)
minAcc = sc.accumulator(np.inf, minAccum())
rdd.foreach(lambda i: minAcc.add(i))
print(minAcc.value)

```




Spark System: Review

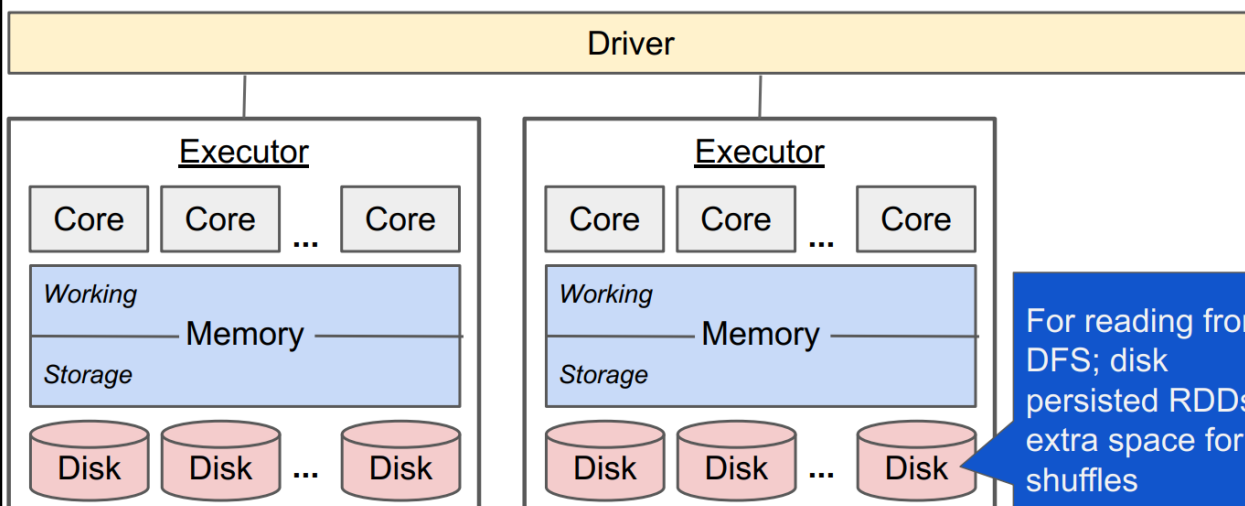
- RDD provides full recovery by backing up transformations from stable storage rather than backing up the data itself.
- RDDs, which are immutable, can be stored in memory and thus are often much faster.
- Functional programming is used to define transformation and actions on RDDs.



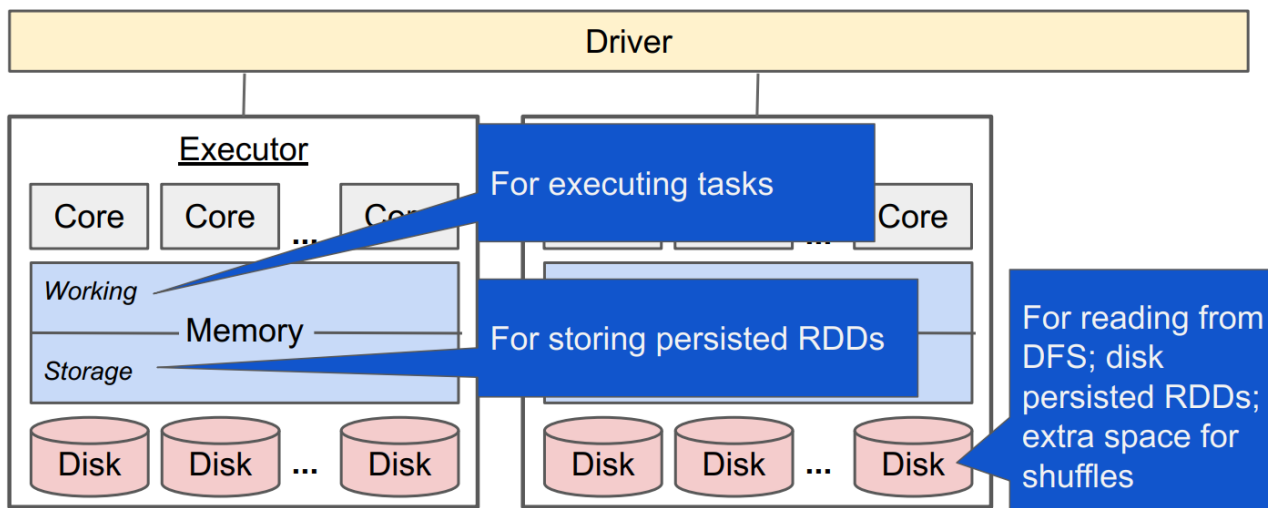
Spark System: Hierarchy



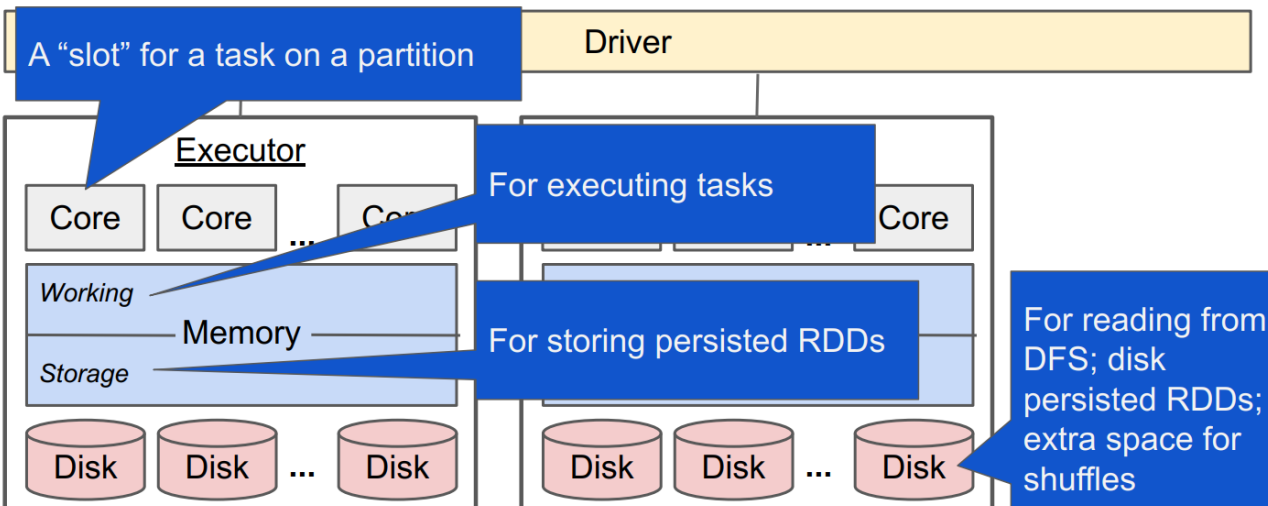
Spark System: Hierarchy



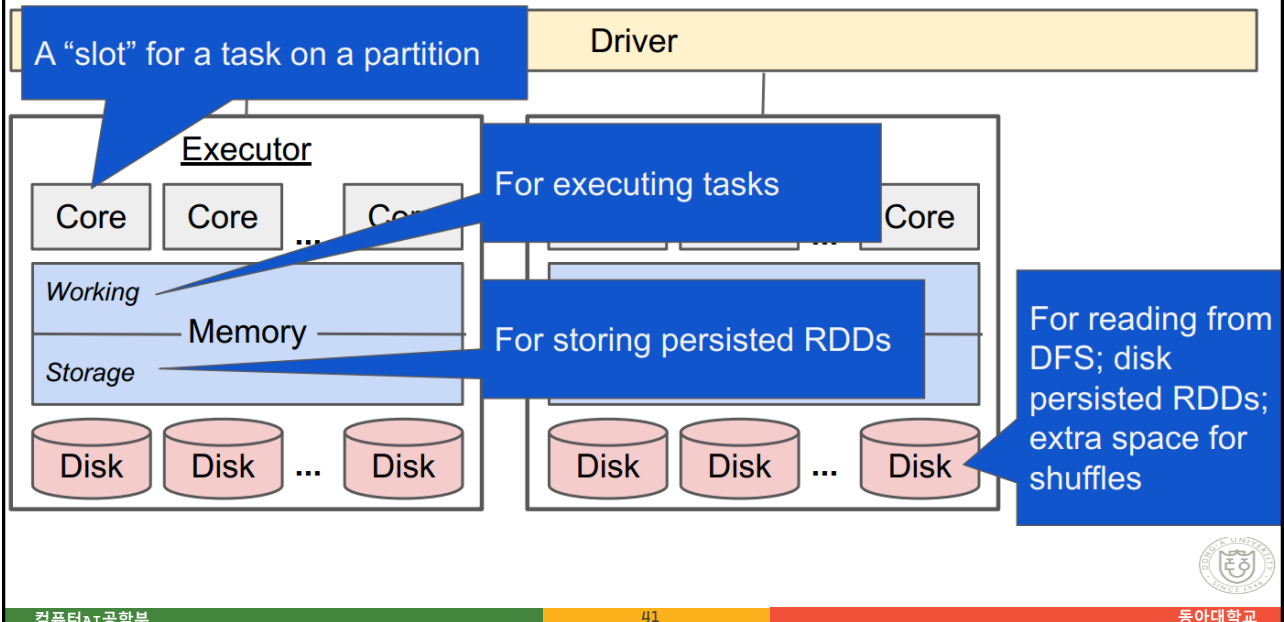
Spark System: Hierarchy



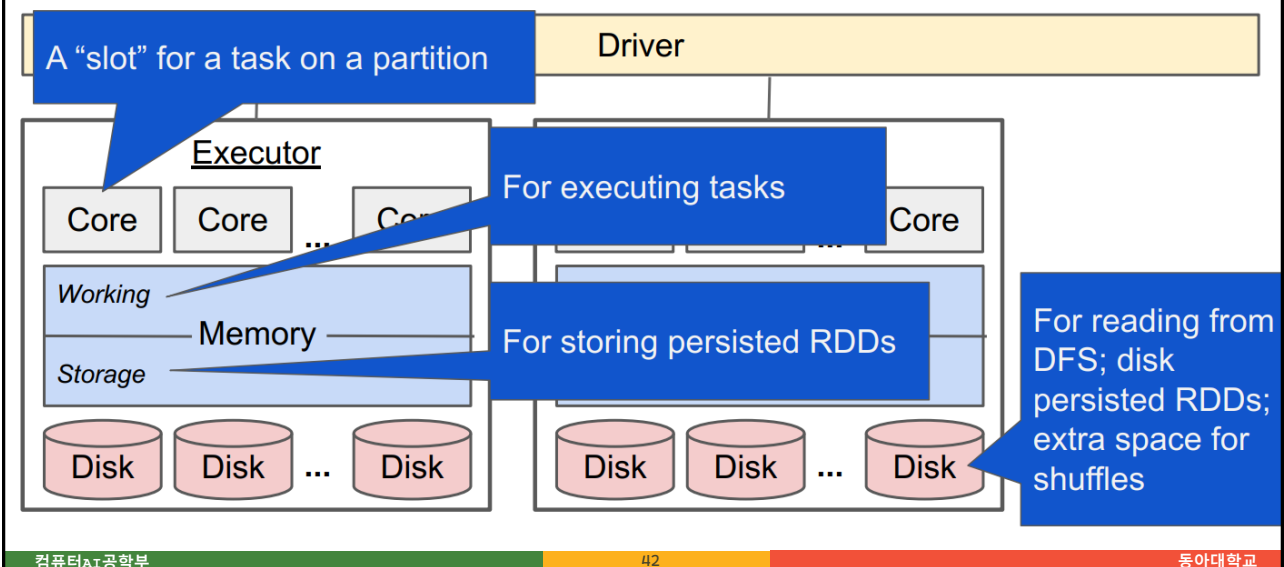
Spark System: Hierarchy



Spark System: Hierarchy



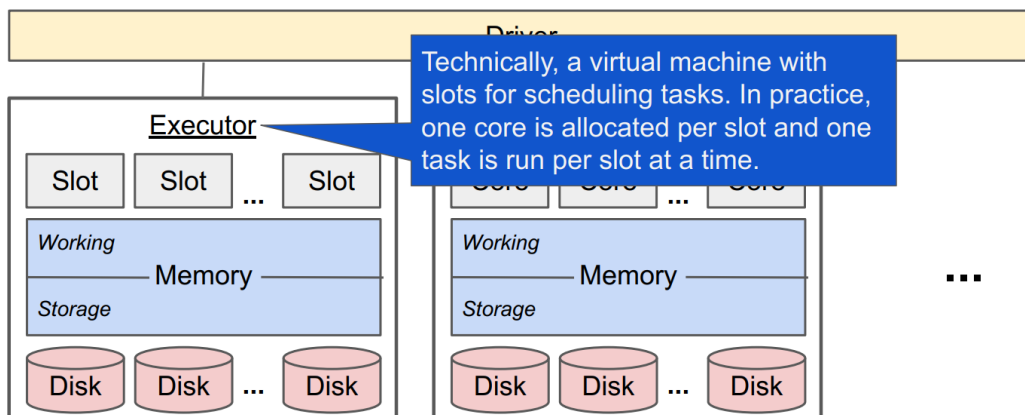
Eager action -> sets off (lazy) chain of *transformations*
 -> launches *jobs* -> broken into *stages* -> broken into **tasks**



Spark System: Hierarchy

Eager action -> sets off (lazy) chain of *transformations*

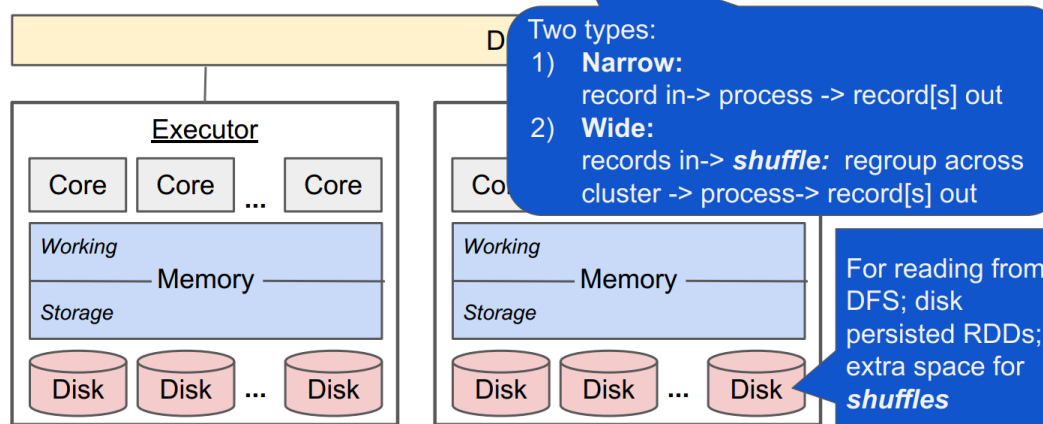
-> launches *jobs* -> broken into *stages* -> broken into *tasks*



Spark System: Hierarchy

Eager action -> sets off (lazy) chain of *transformations*

-> launches *jobs* -> broken into *stages* -> broken into *tasks*



Spark System: Hierarchy

Eager action -> sets off (lazy) chain of **transformations**

-> launches *jobs* -> broken into *stages* -> broken into *tasks*




Image from Nguyen: <https://trongkhoanguyen.com/spark/understand-rdd-operations-transformations-and-actions/>



Spark System: Hierarchy

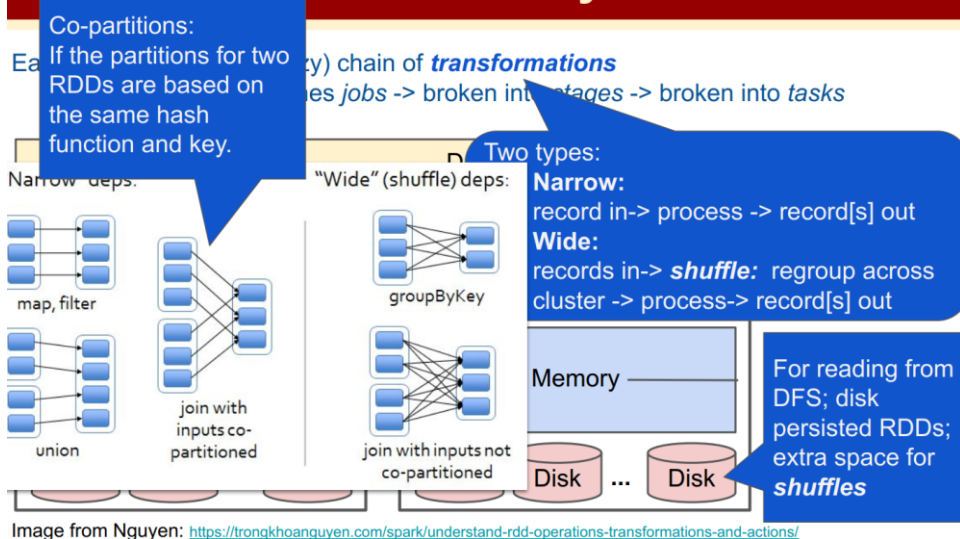


Image from Nguyen: <https://trongkhoanguyen.com/spark/understand-rdd-operations-transformations-and-actions/>




Spark System: Hierarchy

Eager action -> sets off (lazy) chain of transformations
 -> launches jobs -> broken into stages -> broken into tasks

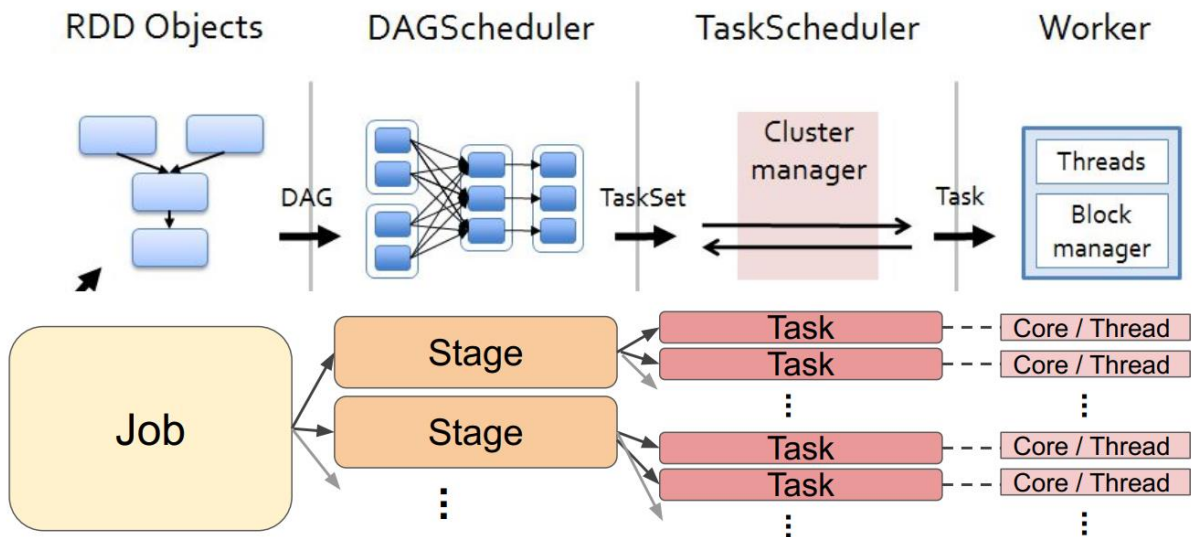
Jobs: A series of transformations (in a DAG) needed for the action

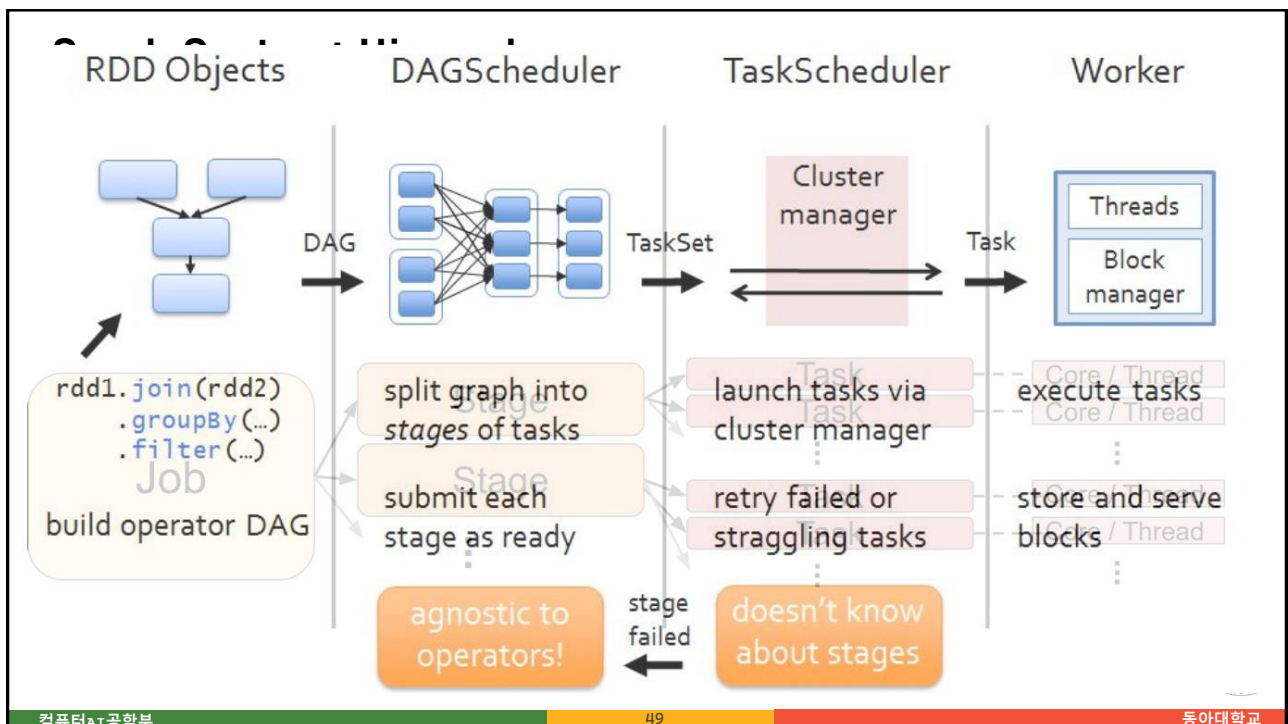
Stages: 1 or more per job -- 1 per set of operations separated by shuffle

Tasks: many per stage -- repeats exact same operation per partition



Spark System: Hierarchy





MapReduce or Spark?

- Spark is typically faster
 - RDDs in memory
 - Lazy evaluation enables optimizing chain of operations.
- Spark is typically more flexible (custom chains of transformations)



MapReduce or Spark?

- Spark is typically faster
 - RDDs in memory
 - Lazy evaluation enables optimizing chain of operations.
- Spark is typically more flexible (custom chains of transformations)

However:

- Still need Hadoop (or some DFS) to hold original or resulting data efficiently and reliably.
- Memory across Spark cluster should be large enough to hold entire dataset to fully leverage speed.

Thus, MapReduce may sometimes be more cost-effective for very large data that does not fit in memory.

